

Data Wrangling in R with the Tidyverse (Part 2)

More data wrangling



Changing *all* columns at once

`mutate_all()`, `rename_all()` : applies a function to *all* columns:

`mutate_all(FUNCTION, FUNCTION_ARGUMENTS)`

```
# mutate_all changes the data in all columns  demo_data %>% mutate_all(as.character) %>%
head(2)
```

1	931897 15 years old Female 10th White	17.179	54.43
2	333862 17 years old Female 12th White	20.2487	57.15
#	rename_all changes all column names		
# A tibble: 3 x 8 demo_data %>% rename_all(str_sub, end = 2) %>% head(3)			

	ag	se	gr	ra	ra	bm	st
	<dbl>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>
1	931897	15	years old	Female	10th	White	White
2	333862	17	years old	Female	12th	White	White
3	36253	18	years old or...	Male	11th	Hispanic/La...	Hispanic/L...

Changing *some* columns at once

`mutate_at()`, `rename_at()`: uses `vars()` to select specific variables to apply a function to i.e.

`mutate_at(vars(SELECT), FUNCTION, FUNCTION_ARGUMENTS)`

```
# mutate_at changes the data in specified columns
demo_data %>% mutate_at(vars(contains("race"), sex), as.factor) %>% head(2)
```

```
# A tibble: 2 x 8
  record age      sex    grade race4 race7    bmi stweight
  <dbl> <chr>    <fct> <chr> <fct> <fct> <dbl>    <dbl>
1  931897 15 years old Female 10th  White White  17.2    54.4
2  333862 17 years old Female 12th  White White  20.2    57.2
```

```
# rename_at changes specified column names
demo_data %>% rename_at(vars(record:grade), toupper) %>% head(3)
```

```
# A tibble: 3 x 8
  RECORD AGE      SEX    GRADE race4      race7    bmi stweight
  <dbl> <chr>    <chr> <chr> <chr>    <chr>    <dbl>    <dbl>
1  931897 15 years old Female 10th  White    White    17.2    54.4
2  333862 17 years old Female 12th  White    White    20.2    57.2
3   36253 18 years old o... Male   11th  Hispanic/L... Hispanic/...  NA      NA
```

Changing *some* columns at once

`mutate_if()`, `rename_if()`, `select_if()`: uses a function that returns TRUE/FALSE to select columns and applies function on the TRUE columns:

`mutate_if(BOOLEAN, FUNCTION, FUNCTION_ARGUMENTS)`

```
demo_data %>% mutate_if(is.numeric, round, digits = 0) %>% head(3)
```

# A tibble: 3 x 8						
record age	sex	grade	race4	race7	bmi	stweight
<dbl> <chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	

demo_data %>% rename_if(is.character,		str_sub,	end = 2)	%>%	head(3)
# A tibble: 3 x 8					
record age	sex	grade	race4	race7	bmi stweight
<dbl> <chr>	<chr>	<chr>	<chr>	<chr>	<dbl> <dbl>

Add one or more rows: `add_row()`

```
arrange(record) %>% head(3)
```

A tibble: 3 x 8

	record	age	sex	grade	race4	race7	bmi	stweight
	<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>
1	100	<NA>	Female	9th	<NA>	<NA>	NA	NA
2	30592	<NA>	<NA>	<NA>	<NA>	<NA>	NA	NA
3	30593	<NA>	<NA>	9th	Hispanic/Latino	Hispanic/Latino	NA	NA

```
demo_data %>% add_row(record=100:102, bmi=c(25,30,18)) %>%  
arrange(record) %>% head(3)
```

A tibble: 3 x 8

record	age	sex	grade	race4	race7	bmi	stweight
<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>

1	100	<NA>	<NA>	<NA>	<NA>	<NA>	25	NA
2	101	<NA>	<NA>	<NA>	<NA>	<NA>	30	NA
3	102	<NA>	<NA>	<NA>	<NA>	<NA>	18	NA

Add one or more columns: `add_column()`

```
head(3)
```

```
# A tibble: 3 x 9
  record study_date age      sex  grade race4      race7    bmi stweight
<dbl> <chr>      <chr>    <chr> <chr> <chr> 1    <dbl>    <dbl>
1 931897 2019-04-10 15 years ... Female 10th White White 17.2    54.4
2 333862 2019-04-10 17 years ... Female 12th White White 20.2    57.2
3 36253 2019-04-10 18 years ... Male 11th Hispani... Hispan... NA      NA
```

```
head(3)
```

```
# A tibble: 3 x 9
   id record age      sex  grade race4      race7    bmi stweight
<int> <dbl> <chr>    <chr> <chr> <chr> 1    <dbl>    <dbl>
1     1 931897 15 years old Female 10th White White 17.2    54.4
2     2 333862 17 years old Female 12th White White 20.2    57.2
3     3 36253 18 years ol... Male 11th Hispanic... Hispanic... NA      NA
```

Frequency tables: janitor package's tabyl function

```
# default table
demo_data %>% tabyl(grade)
```

```
grade    n percent valid_percent
10th  4907 0.24535      0.2504338
11th  4891 0.24455      0.2496172
12th  4577 0.22885      0.2335919
9th   5219 0.26095      0.2663570
<NA>   406 0.02030           NA
```

```
# output can be treated as tibble
demo_data %>% tabyl(grade) %>% select(-n)
```

```
grade percent valid_percent 10th 0.24535
      0.2504338
11th 0.24455      0.2496172
12th 0.22885      0.2335919
9th  0.26095      0.2663570
<NA> 0.02030           NA
```

adorn_ your table!


```
demo_data %>% tabyl(grade) %>%
```

```
  adorn_totals("row") %>%
  adorn_pct_formatting(digits=2)
```

```
grade          n percent
valid_percent 10th      4907 25.04%
              24.54%      24.96%
11th      4891  24.45%      23.36%
12th      4577  22.88%      26.64%
9th       5219  26.10%      -
Total 20406 100.00%      100.00%
```


2x2 tabyls

# default 2x2 table demo_data %>% tabyl(grade, sex)			
grad	Femal	Male NA_	
12th	2277	2263 37	
9th	2492	2684 43	
<NA>	126	195 85	



grade	Female	Male	NA_
-------	--------	------	-----

<NA>	1.3% (126)	1.9% (195)	36.8% (85)	
9th	100.0% (2492)	100.0% (2684)	100.0% (43)	
12th	100.0% (2277)	100.0% (2263)	100.0% (37)	

What adornments does the tabyl to right have?

- Notice grade is not sorted in a pleasing way. We will learn how to deal with this when we discuss factors as a data type in R.
- Base R has a table function, but it is clunkier and the output is not a data frame. See the [tabyl vignette](#) for more information, adorn options, & 3-way tabyls

Numerical data summaries: `summarize()`

- We can summarize data as a whole, or in groups with `group_by()`
- `group_by()` is very powerful, see [data wrangling cheatsheet](#)
- Can also use `summarize_at()`, `summarize_if()`, `summarize_all()`

```
# summary of all data as a whole demo_data
%>%
  summarize(bmi_mean = mean(bmi, na.rm=TRUE),
            bmi_sd = sd(bmi, na.rm=TRUE))
```

```
# A tibble: 1 x 2 bmi_mean bmi_sd
      <dbl> <dbl>
1  23.5    4.99
```

What does `na.rm=TRUE` do and what happens if we leave it out?


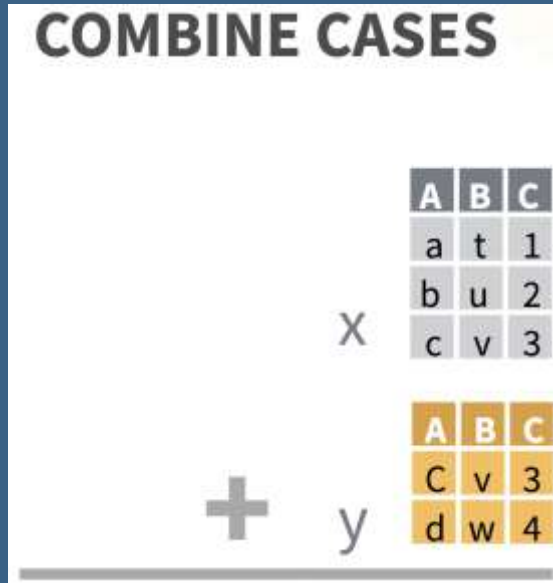
```
# summary by group variable demo_data %>%

summarize(n_per_group = n(),
          bmi_mean = mean(bmi, na.rm=TRUE),
          bmi_sd = sd(bmi, na.rm=TRUE))
```

```
# A tibble: 5 x 4
  grade n_per_group bmi_mean bmi_sd
  <chr>      <int>      <dbl> <dbl>
```

1 <NA>	406	23.5	6.45
2 10th	4907	23.2	4.76
3 11th	4891	23.8	4.89
4 12th	4577	24.2	5.20
5 9th	5210	22.8	4.82

Combining data sets



Rows (cases): paste data below each other

`bind_rows()` combines rows from different data sets
& accounts for different column names

```
data1
```

```
bind_rows(data1,data2, .id = "group")
```

#	A tibble:	x 4			#	A tibble:	x 6		
	2					5			
	id name	height	age			group id	name	height	age years
	<int> <chr>	<dbl>	<dbl>			<chr> <int>	<chr>	<dbl>	<dbl> <dbl>
1	1 Nina	2	4		1	1 Nina	2	4	NA
2	2 Yi	1	2		2				NA
2	8 Al	1.7	1		3	2	2	1	3
3	9 Juan	1.8	2			7			

DF	A	B	C
x	a	t	1
x	b	u	2
x	c	v	3
z	c	v	3
z	d	w	4

bind_rows(..., .id = NULL)
Returns tables one on top of the other as a single table. Set .id to a column name to add a column of the original table names (as pictured)

Columns (variables): *DO NOT USE* `bind_cols()` !!

- `bind_cols()` blindly pastes columns next to each other without preserving order of variables that they have in common
 - Use `join` to preserve ordering - see next slides

# datasets must have same number of rows to use bind_cols()											
demo_sub <- demo_data %>% slice(1:20) qn_sub <- qn_data %>% slice(1:20)						# first 20 rows of demo_data # first 20 rows of qn_data					
bind_cols(demo_sub, qn_sub)						# blindly bind columns; 2nd record column got renamed					

joining your data sets

- `Join` uses overlapping or selected columns to combine two or more data sets.
- Also called "merging" or "mutating join".
- Function names are based off of SQL operations for databases.

X						y					
A	B	C				A	B	D			
a	t	1				a	t	3			
b	u	2				b	u	2			
c	v	3				d	w	1			

Use a "**Mutating Join**" to join one table to columns from another, matching values with the rows that they correspond to. Each join retains a different combination of values from the tables.

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA

left_join(x, y, by = NULL, copy=FALSE, suffix=c(".x",".y"),...) Join matching values from y to x.

A	B	C	D
a	t	1	3
b	u	2	2
d	w	NA	1

right_join(x, y, by = NULL, copy = FALSE, suffix=c(".x",".y"),...) Join matching values from x to y.


A	B	C	D
a	t	1	3
b	u	2	2

inner_join(x, y, by = NULL, copy = FALSE, suffix=c(".x",".y"),...) Join data. Retain only rows with matches.

A	B	C	D
a	t	1	3
b	u	2	2
c	v	3	NA
d	w	NA	1

full_join(x, y, by = NULL, copy=FALSE, suffix=c(".x",".y"),...) Join data. Retain all values, all rows.

join options visually



[Hiroaki Yutani](#)

Most commonly used: `left_join()`

- `left_join(x, y)` includes all observations in `x`, regardless of whether they match ones in `y` or not.
- It includes all columns in `y`, but only rows that match `x`'s observations.

```
df1 <- tibble(a = c(1, 2), b = 2:1)
df2 <- tibble(a = c(1, 3), c = 10:11)
df1
```

#	A	tibble:	x	2
		2		
		a	b	

```
2      2      1
# A tibble: 2 x 2
      a      c
<dbl> <int>
1     1    10
2     3    11
```

```
left_join(df1, df2)
```

```
# A tibble: 2 x 3
      a      b      c
<dbl> <int> <int>
```

1	1	2	10
2	2	1	NA

- Which common column(s) were used to merge the datasets?
- What if we want to specify which columns to join by when merging? see next slide...

Which columns will be used to join?

- If no columns are specified to join by, then *all* overlapping (intersecting) column names will be used
- Often we want to specify which columns to use,
 - and also how to rename duplicated columns that were not merged

X						y					
A	B	C				A	B	D			
a	t	1				a	t	3			
b	u	2				b	u	2			
c	v	3				d	w	1			

A	B.x	C	B.y	D
a	t	1	t	3
b	u	2	u	2
c	v	3	NA	NA

Use **by = c("col1", "col2", ...)** to specify one or more common columns to match on.
`left_join(x, y, by = "A")`

A.x	B.x	C	A.y	B.y
a	t	1	d	w
b	u	2	b	u
c	v	3	a	t

Use a named vector, **by = c("col1" = "col2")**, to match on columns that have different names in each table.
`left_join(x, y, by = c("C" = "D"))`

A1	B1	C	A2	B2
a	t	1	d	w
b	u	2	b	u
c	v	3	a	t

Use **suffix** to specify the suffix to give to unmatched columns that have the same name in both tables.
`left_join(x, y, by = c("C" = "D"), suffix = c("1", "2"))`

[dplyr data transformation cheatsheet](#)

Check for overlapping column names

Goal: merge the demographics (`demo_data`) and questionnaire (`qn_data`) together.
What column names do these datasets have in common?

colnames(demo_data)			
[1] "record" "age" "sex" [7] "bmi" "stweight"		"grade" "race4" "race7"	
colnames(qn_data)			
[1] "record" "q8" "q12"	"q3"	"qn24"	
intersect(colnames(demo_data), colnames(qn_data))			

```
[1] "record"
```

Merge demo_data and qn_data together

Let's do a full join so that we keep all data from both datasets

```
merged_data <- full_join(demo_data,
  qn_data,
  by = "record")

# Check dimensions of original and new data
```

```
dim(demo_data); dim(qn_data); dim(merged_da
```

```
[1] 20000      8
```

```
[1] 10000      5
```

```
[1] 20000     12
```

```
merged_data
```


#	A tibble: 20,000 x 12												
	record	age	sex	grade	race4	race7	bmi	stweight	q8	q12	q31		
	<dbl>	<chr>	<chr>	<chr>	<chr>	<chr>	<dbl>	<dbl>	<chr>	<chr>	<chr>		
	1	9.32e5	15	y...	Fema...	10th	White	White	17.2	54.4	Nev e...	<NA>	Yes
	2	3.34e5	17	y...	Fema...	12th	White	White	20.2	57.2	<NA>	<NA>	<NA>

Practice

1. Add a column of 1's to `qn_data` called `qn_yes` and save the resulting data as `qn_data2`.
2. Join `demo_data` and `qn_data2` by column `record`. Keep all rows from `demo_data` and only rows from `qn_data2` that match records in `demo_data`. Call the resulting data `all_data`.
3. Create a `tabyl()` of `qn_yes` for the data `all_data`.
4. Create a 2x2 table of `qn_yes` vs `grade`. Note about the data:
 - q8 = How often wear bicycle helmet
 - q12 = Texted while driving
 - q31 = Ever smoked
 - qn24 = Bullied past 12 months
 -
 -
 -

Reshaping data

wide vs. long



[Allison Horst](#)



[tidyr](#)

Wide vs. long data

- **Wide** data has one row per subject, with multiple columns for their repeated measurements
- **Long** data has multiple rows per subject, with one column for the measurement variable and another indicating from when/where the repeated measures are from

wide

id	SBP_visit1	SBP_visit2	SBP_visit3
a	130	110	112
b	120	116	122
c	130	136	138
d	119	106	118

long

id	visit	SBP
a	1	130
b	1	120
c	1	130
d	1	119
a	2	110
b	2	116
c	2	136
d	2	106
a	3	112
b	3	122
c	3	138
d	3	118

Example wide dataset

Copy and paste the code below into R to create this example dataset

```
BP_wide <- tibble(id = letters[1:4],  
                  sex = c("F", "M", "M", "F"), SBP_v1 = c(130, 120, 130, 119),  
                  SBP_v2 = c(110, 116, 136, 106),  
                  SBP_v3 = c(112, 122, 138, 118))  
  
BP_wide
```

# A tibble:		x 5		
4				
	id	sex	SBP_v1	SBP_v2 SBP_v3
	<chr>		<dbl>	<dbl> <dbl>
	<chr>			
1	a	F	130	110 112
2	b	M	120	116 122
3	c	M	130	136 138
4	d	F	119	106 118

- What do you think the data in the table are measures of?
- How can we tell the data are wide?

Wide to long: gather()

```
BP_wide
```

```
# A tibble: 4 x 5
  id    sex  SBP_v1 SBP_v2 SBP_v3
  <chr> <chr> <dbl>   <dbl>
1 a      F      130     110
2 b      M      120     116     122
3 c      M      130     136     138
4 d      F      119     106     118
```

gather columns into rows to make the data long. Need to **specify**:

- **new column names**
 - **key**: stores row names of wide data's gathered columns
 - **value**: stores data values
- **which columns to gather**

```
BP_long <- BP_wide %>%
  gather(key = "visit", value = "SBP",
         SBP_v1:SBP_v3)
```

```
BP_long
```

```
# A tibble: 12 x 4
  id    sex  visit      SBP
  <chr> <chr> <chr>   <dbl>
1 a      F    SBP_v1    130
2 b      M    SBP_v1    120
```

3	c	M	SBP_v1	130
4	d	F	SBP_v1	119
5	a	F	SBP_v2	110
6	b	M	SBP_v2	116
7	c	M	SBP_v2	136
8	d	F	SBP_v2	106
9	a	F	SBP_v3	112
10	b	M	SBP_v3	122
11	c	M	SBP_v3	138

Long to wide: spread()

BP_long

```
# A tibble: 12 x 4
  id    sex  visit    SBP
<chr> <chr> <chr>   <dbl>
1 a      F    SBP_v1    130
2 b      M    SBP_v1    120
3 c      M    SBP_v1    130
4 d      F    SBP_v1    119
5 a      F    SBP_v2    110
6 b      M    SBP_v2    116
7 c      M    SBP_v2    136
8 d      F    SBP_v2    106
9 a      F    SBP_v3    112
10 b     M    SBP_v3    122
11 c     M    SBP_v3    138
12 d     F    SBP_v3    118
```

spread rows into columns to make the data wide. Need to **specify** which columns in the long data to use:

- **key** column: has the variable names
- **value** column: has the data values

```
BP_wide2 <- BP_long %>%
  spread(key = "visit", value = "SBP")
BP_wide2
```

```
# A tibble: 4 x 5
  id    sex  SBP_v1 SBP_v2 SBP_v3
  <chr> <chr> <dbl>   <dbl>
1 a      F    130     110
  112
2 b     M    120     116     122
3 c     M    130     136     138
4 d     F    119     106     118
```

Clean up long data's visit column (key column)

BP_long

```
# A tibble: 12 x 4
```

id		sex	SBP		visit		SBP	
							<dbl>	
1	a	F	SBP_v	130	1	SBP_v	130	1
2	b	M	SBP_v	120	1	SBP_v	120	1
3	c	M	SBP_v1	130	1	SBP_v1	130	1
4	d	F	SBP_v1	119	1	SBP_v1	119	1
5	a	F	SBP_v2	110	2	SBP_v2	110	2
6	b	M	SBP_v2	116	3	SBP_v2	116	3
7	c	M	SBP_v2	136	4	SBP_v2	136	4
8	d	F	SBP_v2	106	5	SBP_v2	106	5
9	a	F	SBP_v3	112	11	SBP_v3	112	11
10	b	M	SBP_v3	122	12	SBP_v3	122	12
11	c	M	SBP_v3	138	8	SBP_v3	138	8

```
BP_long2 <- BP_long %>% mutate(visit =  
  str_replace(visit, "SBP_v", ""))
```

BP_long2

```
# A tibble:
```

```
12
```

```
id
```

```
sex
```

```
x 4
```

```
vis
```

```
it
```

```
SBP
```

```
<ch
```

```
r>
```

```
<ch
```

```
r>
```

```
<ch
```

```
r>
```

```
<dbl
```

```
>
```

```
1
```

```
a
```

```
F
```

```
1
```

```
130
```

```
2
```

```
b
```

```
M
```

```
1
```

```
120
```

```
3
```

```
c
```

```
M
```

```
1
```

```
130
```

```
4
```

```
d
```

```
F
```

```
1
```

```
119
```

```
5
```

```
a
```

```
F
```

```
2
```

```
110
```

```
11
```

```
c
```

```
M
```

```
3
```

```
138
```

```
12
```

```
d
```

```
F
```

```
3
```

```
118
```

```
8
```

```
d
```

```
F
```

```
2
```

```
106
```

Make cleaned-up long data wide

```
head(BP_long2, 2)
```

```
# A tibble: 2 x 4
  id    sex visit  SBP
  <chr> <chr> <chr>
<dbl> 1 a    F      1
      130
2 b M    1      120
```

```
BP_wide3 <- BP_long2 %>%
  spread(key = "visit", value = "SBP")
BP_wide3
```

```
# A tibble: 4 x 5
  id    sex `1`    `2`    `3`
  <chr> <chr> <dbl> <dbl>
<dbl> 1 a    F      130
      110  112
2 b M    120  116  122
4 d 3 c M    130  106  138
```

Problem: have numbers as column names, since `spread`'s default is to use the levels of the `key` as the new row names.

Solution: have row names start with the `key` column's name separated by a character

```
BP_wide4 <- BP_long2 %>%
  spread(key = "visit", value = "SBP",
        sep="_") # specify separating char
BP_wide4
```

```
# A tibble: 4 x 5
  id    sex visit_1 visit_2 visit_3
  <chr> <chr> <dbl>    <dbl>    <dbl>
1 a    F      130
      110  112
2 b M    120  116  122
4 d 3 c M    130  106  138
```

Practice

Copy and paste the code below into R to create the dataset `DBP_wide`

```
DBP_wide <- tibble(id = letters[1:4],  
                  sex = c("F", "M", "M", "F"), v1.DBP = c(88, 84, 102, 70),  
                  v2.DBP = c(78, 78, 96, 76),  
                  v3.DBP = c(94, 82, 94, 74),  
                  age=c(23, 56, 41, 38)  
                  )
```

1. Make `DBP_wide` into a long dataframe based on the repeated DBP columns and save it as `DBP_long`.
2. Clean up the visit column of `DBP_long` so that the values are 1,2,3, and save it as `DBP_long`.
3. Make `DBP_long` wide with column names `visit.1`, `visit.2`, `visit.3` for the DBP values, and save it as `DBP_wide2`.
4. Join `DBP_long` with `BP_long2` so that we have one data frame with columns `id`, `sex`, `visit`, `SBP`, `DBP`, and `age`. Save this as `BP_both_long`.

Data cleaning

(messy NAs, names, strings, dates, factors)



Removing missing data: drop_na()

A small data example:

```
mydata <- tibble(id = 7:9,  
                 name = c("Bo", "Al", "Juan"),  
                 height = c(2, NA, 1.8),  
                 years = c(51, 35, NA))  
  
mydata
```

```
# A tibble: 3 x 4  
   id name  height years  
  <int> <chr>   <dbl> <dbl>
```

1	7	Bo	2	51
2	8	Al	NA	35
3	9	Juan	1.8	NA

Remove *all* rows with **any missing data**

```
mydata %>% drop_na()
```

```
# A tibble: 1 x 4  
       id name height years  
   <int> <chr>   <dbl> <dbl> 1     7 Bo  
                                     2     51
```

Remove rows with NA in **selected columns**

```
mydata %>% drop_na(height)
```

```
# A tibble: 2 x 4
```

	id	name	height	years
1	7	Bo	2	51
2	9	Juan	1.8	NA

Replace NAs with another value: `replace_na()`

Use with `mutate()`

mydata

```
# A tibble: 3 x 4  
  id name height years  
  <int> <chr>   <dbl> <dbl>
```

1	7 Bo	2	51
2	8 Al	NA	35
3	9 Juan	1.8	NA

```
mydata %>%
```

```
  mutate(height = replace_na(height, "Unknown"),  
         years = replace_na(years, 0) )
```

# A tibble:		x 4	
3			
id name		height	
2	8 Al	Unknown	35
3	9 Juan	1.8	0
		1>	
1	7 Bo	2	51

replace_na() advanced example

Replaces `NA`s in all columns starting with "q" with the string "No answer"

```
qn_data %>%
```

```
  tabyl(q8, q31)
```

q8	No	No o	answe r	Yes
Always wore a helmet	411		58	216
Did not ride a bicycle	131 8		343	122 3
Most of the time wore a helmet	320		31	173
Never wore a helmet	142 9		418	167 2
No answer	267		212	300
Rarely wore a helmet	481		73	405
Sometimes wore a helmet	351		53	246

Convert (i.e. "No answer", 9999, etc) to NA: `na_if()`

```
all_data %>% tabyl(race4)
```

race4	n percent	valid_percent
All other races	4713 0.23565	0.2443235
Black or African American	4093 0.20465	0.2121825

```
all_data %>%
```

```
  tabyl(race4)
```

		0.29070	
<NA>	710 0.03550		NA
race4	n percent	valid_percent	Black or African American
	4093 0.20465		0.2807848
Hispanic/Latino	4670 0.23350		0.3203677
White	5814 0.29070		0.3988475
<NA>	5423 0.27115		NA

na_if() for all your data

Avoid this by reading in your data correctly:

```
smalldata <- read_csv("data/small_data.csv",  
                      na = c("", "9999", "NA")) # specify your own missing values
```

Otherwise na_if() everything:

```
# replace all "" with NA all_data %>%  
  
# replace all 9999's with NA all_data %>%  
  
mutate_if(is.numeric, .funs = na_if(., 9999))
```

Working with character strings

- Use the [package stringr](#) (loaded with `tidyverse`)
- Paste strings or values together [with package glue](#) (installed, not loaded w/ `tidyverse`)
- *advanced tip*: learn ["regular expressions"](#) ([regex](#)) for pattern matching (see [cheatsheet](#)) and matching multiple characters/strings at once

String manipulation with stringr : : CHEAT SHEET

The **stringr** package provides a set of internally consistent tools for working with character strings, i.e. sequences of characters surrounded by quotation marks.



Detect Matches



str_detect(string, pattern) Detect the presence of a pattern match in a string.
`str_detect(fruit, "a")`



str_which(string, pattern) Find the indexes of strings that contain a pattern match.
`str_which(fruit, "a")`




str_count(string, pattern) Count the number of matches in a string.
`str_count(fruit, "a")`




str_locate(string, pattern) Locate the positions of pattern matches in a string. Also **str_locate_all**. `str_locate(fruit, "a")`


Subset Strings



str_sub(string, start = 1L, end = -1L) Extract substrings from a character vector.
`str_sub(fruit, 1, 3); str_sub(fruit, -2)`



str_subset(string, pattern) Return only the strings that contain a pattern match.
`str_subset(fruit, "b")`



str_extract(string, pattern) Return the first pattern match found in each string, as a vector. Also **str_extract_all** to return every pattern match. `str_extract(fruit, "[aeiou]")`



str_match(string, pattern) Return the first pattern match found in each string, as a matrix with a column for each () group in pattern. Also **str_match_all**.
`str_match(sentences, "[a]the ([^]+)")`

Manage Lengths



str_length(string) The width of strings (i.e. number of code points, which generally equals the number of characters). `str_length(fruit)`



str_pad(string, width, side = c("left", "right", "both"), pad = " ") Pad strings to constant width. `str_pad(fruit, 17)`



str_trunc(string, width, side = c("right", "left", "center"), ellipsis = "...") Truncate the width of strings, replacing content with ellipsis. `str_trunc(fruit, 3)`



str_trim(string, side = c("both", "left", "right")) Trim whitespace from the start and/or end of a string. `str_trim(fruit)`

[stringr cheatsheet](#)

str_detect() find strings

```
mydata <- tibble(name = c("J.M.", "Ella", "Jay"), state = c("New Mexico", "New York", "Oregon"))
```

Filter based on string detection

```
mydata %>% filter(str_detect(name, "J"))
```

```
# A tibble: 2 x 2 name state
<chr> <chr>
1 J.M. New Mexico
2 Jay Oregon
```

Creates a column of TRUE/FALSE if detected

```
mydata %>% mutate(
  new_state = str_detect(state, "New"))
```

```
# A tibble: 3 x 3
  name state new_state
<chr> <chr> <lgl>
1 J.M. New Mexico TRUE
2 Ella New York TRUE
3 Jay Oregon FALSE
```

str_replace_all(), str_replace()

```
mydata %>% mutate(state_old = str_replace_all(state, "New", "Old"))
```

```
# A tibble: 3 x 3
  name    state    state_old
<chr> <chr>    <chr>
1 J.M.   New Mexico Old Mexico
2 Ella   New York   Old York
3 Jay    Oregon     Oregon
```

```
mydata %>% mutate(
  name2 = str_replace(name, "l", "-"),
  name3 = str_replace_all(name, "l", "-"),
  name4 = str_replace_all(name,
    fixed("."),
    fixed("."))
)
```

name2 = str_replace(name, "l", "-"),		# first instance		
name3 = str_replace_all(name, "l", "-"),		# all instances		
name4 = str_replace_all(name, fixed("."), fixed("."))	"")	# special characters	wi th	fixed ()

```
# A tibble: 3 x 5
  name    state    name2 name3 name4
<chr> <chr>    <chr> <chr> <chr>
1 J.M.   New Mexico J.M.   J.M.   JM
2 Ella   New York   E-la   E--a   Ella
3 Jay    Oregon     Jay    Jay    Jay
```

str_sub() : shorten strings

Based on position 1 (start = 1) to length of string (end = -1)

mydata %>% mutate(

short_name = str_sub(name, start = 1, end = 2),	#	specify start to end
short_name2 = str_sub(name, end = 2),	#	specify only end
short_state = str_sub(state, end = -3)	#	negative endices, from end
)		

#	A tibble: 3 x 5		
	name state	short_name short_name2	short_state
	<chr> <chr>	<chr> <chr>	<chr>
1	J.M. New Mexico	J.	New Mexi
2	Ella New York	El	New Yo
3	Jay Oregon	Ja	Oreg

Paste strings together with `glue()`

- `paste()` is the base R way of pasting strings (surprise, it's hard to use)
- `glue()` is most useful when pasting data columns together
- **column names or function operations go inside `{ }`**
- See the [glue vignette](#)

```
all_data %>%
```

```
  select(record, info) %>% head(5)
```

```
# A tibble: 5 x 2 record info
```

```
<dbl> <S3: glue>
```

```
1  931897 Student 931897 is 15 years old with BMI = 17.2
2  333862 Student 333862 is 17 years old with BMI = 20.2
3   36253 Student 36253 is 18 years old or older with BMI = NA
4 1095530 Student 1095530 is 15 years old with BMI = 28
5 1303997 Student 1303997 is 14 years old with BMI = 24.5
```

Using glue to summarize data

- Useful for tables (will cover this more in another session)
- Example, calculate the S.E. of the mean and create a column with "mean (SE)" of bmi:

```
demo_data %>% group_by(sex) %>% summarize(n_sex = n(),  
      bmi_mean = mean(bmi, na.rm=TRUE), bmi_sd = sd(bmi, na.rm=TRUE)) %>%
```

```
mutate(bmi_mean_se = glue("{round(bmi_mean,1)} ({signif(bmi_sd/sqrt(n_sex),2)})"))
```

```
# A tibble: 3 x 5
```

```
sex      n_sex bmi_mean bmi_sd bmi_mean_se  
<chr>   <int>   <dbl>  <dbl> <S3: glue>
```

1	<NA>	231	NaN	NaN	NaN (NaN)
2	Female	9592	23.3	4.96	23.3 (0.051)
3	Male	10177	23.7	5.01	23.7 (0.05)

Wrangle dates with `lubridate`

- Convert characters to special "Date" type
- Convert *terrible excel date formats* into workable data
- Easy date magic examples: add and
 - subtract datesconvert to minutes/years/etc change
 - timezonesadd 1month to a date...
- [lubridate cheat sheet](#)
- `read_csv` and `read_excel` etc automatically import dates correctly



[Allison Horst](#)

What kind of date do you have?

PARSE DATE-TIMES (Convert strings or numbers to date-times)

1. Identify the order of the year (**y**), month (**m**), day (**d**), hour (**h**), minute (**m**) and second (**s**) elements in your data.
2. Use the function below whose name replicates the order. Each accepts a wide variety of input formats.

2017-11-28T14:02:00 `ymd_hms()`, `ymd_hm()`, `ymd_h()`.
`ymd_hms("2017-11-28T14:02:00")`

2017-22-12 10:00:00 `ydm_hms()`, `ydm_hm()`, `ydm_h()`.
`ydm_hms("2017-22-12 10:00:00")`

11/28/2017 1:02:03 `mdy_hms()`, `mdy_hm()`, `mdy_h()`.
`mdy_hms("11/28/2017 1:02:03")`

1 Jan 2017 23:59:59 `dmy_hms()`, `dmy_hm()`, `dmy_h()`.
`dmy_hms("1 Jan 2017 23:59:59")`

20170131 `ymd()`, `ydm()`. `ymd(20170131)`

July 4th, 2000 `mdy()`, `myd()`. `mdy("July 4th, 2000")`

4th of July '99 `dmy()`, `dym()`. `dmy("4th of July '99")`

2001: Q3 `yq()` Q for quarter. `yq("2001: Q3")`

2:01 `hms::hms()` Also `lubridate::hms()`,
`hm()` and `ms()`, which return
periods.* `hms::hms(sec = 0, min = 1,
hours = 2)`

[lubridate cheat sheet](#)

```
timedata <-  
  tibble(name = c("Yi", "Bo", "DJ"),  
         dob=c("10/31/1952", "1/12/1984", "2/02/2002"))  
timedata %>%  
  mutate(dob_date = mdy(dob),  
         dob_wrong = dmy(dob)) # wrong order
```

# A tibble: 3 x 4			
name	dob	dob_date	dob_wron g
<chr> <chr>		<date>	<date>
1 Yi	10/31/19	1952-10-31	NA
52			
2 Bo	1/12/198	1984-01-12	1984-12-01
4			
3 DJ	2/02/200	2002-02-02	2002-02-02
2			

Math with dates

```
timedata %>% mutate(
```

<code>dob = mdy(dob),</code>	#	convert to a date
<code>dob_year = year(dob),</code>	#	extract the year
<code>time_since_birth = dob %--% today(),</code>	#	create an "interval"
<code>age = time_since_birth %/% years(1),</code>	#	modulus on "years"
<code>dobplus = dob + days(10)</code>	#	add 10 days
# A) tibble: 3 x 6		
<code>name dob</code>	<code>dob_year</code>	<code>time_since_birth</code>

`name dob dob_year time_since_birth age dobplus`
<chr> <date> <dbl> <S4: Interval> <dbl> <date>



1	Yi	1952-10-31	1952 1952-10-31 UTC--2019-05-23	UTC	66 1952-11-10
2	Bo	1984-01-12	1984 1984-01-12 UTC--2019-05-23	UTC	35 1984-01-22
3	DJ	2002-02-02	2002 2002-02-02 UTC--2019-05-23	UTC	17 2002-02-12

Factors - categorical data

- Clean and order factors with `forcats` package
- Will go over this for `ggplot2` (visualization), statistical modeling (i.e. for `lm()`), and probably a workshop for creating tables
- See [forcats cheatsheet](#) and `forcats` [vignette](#)


Factors

R represents categorical data with factors. A **factor** is an integer vector with a **levels** attribute that stores a set of mappings between integers and categorical values. When you view a factor, R displays not the integers, but the values associated with them.



Create a factor with `factor()`

factor(x = character(), levels, labels = levels, exclude = NA, ordered = is.ordered(x), nmax = NA) Convert a vector to a factor. Also **as_factor**.
f <- factor(c("a", "c", "b", "a"), levels = c("a", "b", "c"))



Return its levels with `levels()`

levels(x) Return/set the levels of a factor. *levels(f); levels(f) <- c("x","y","z")*

Use `unclass()` to see its structure

forcats examples - specify levels fct_relevel()

```
mydata <- tibble( id = 1:4,  
  grade=c("9th", "10th", "11th", "9th")) %>%
```

```
levels(mydata$grade_fac)
```

```
[1] "10th" "11th" "9th"
```

```
mydata %>% arrange(grade_fac)
```

```
# A tibble: 4 x 3  
      id grade grade_fac  
  <int> <chr> <fct>  
1     2 10th  10th  
2     3 11th  11th  
3     1  9th   9th  
4     4  9th   9th
```

```
mydata <- mydata %>% mutate(
```

```
levels(mydata$grade_fac)
```

```
[1] "9th"  "10th" "11th"
```

```
mydata %>% arrange(grade_fac)
```

```
# A tibble: 4 x 3  
      id grade grade_fac  
  <int> <chr> <fct>  
1     1  9th   9th  
2     4  9th   9th  
3     2 10th  10th  
4     3 11th  11th
```

forcats examples - collapse levels

```
mydata <- tibble(loc = c("SW", "NW", "NW", "NE", "SE", "SE"))

mydata %>% mutate( loc_fac = factor(loc),

  loc2 = fct_collapse(loc_fac,                                # collapse levels
    south = c("SW", "SE"),
    north = c("NW", "NE"),
    other = c("other") # most common 2 levels + other
  )
```

A tibble: 6 x 4

	loc	loc_fac	loc2	loc3
	<chr>	<fct>	<fct>	<fct>
1	SW	SW	south	other
2	NW	NW	north	NW
3	NW	NW	north	NW
4	NE	NE	north	other
5	SE	SE	south	SE
6	SE	SE	south	SE

Other "janitor" work



Clean messy column names with `clean_names()`

```
mydata <- tibble("First Name"= c("Yi","DJ"), "last init" = c("C","R"),
                 "% in" = c(0.1, 0.5), "ñ$$$$"= 1:2, " "=3:2, " hi"=c("a","b"),
                 "null"=c(NA,NA))

mydata
```

```
# A tibble: 2 x 7
  `First Name` `last init` `% in` `ñ$$$$` ` ` ` hi` null
  <chr>        <chr>      <dbl> <int> <int> <chr> <lg
  <1> 1 Yi      C          0.1    1    3 a    NA
  <2> 2 DJ      R          0.5    2    2 b    NA
```

```
remove_empty(c("rows","cols")) # also useful
```

```
# A tibble: 2 x 6
  first_name last_init percent_in     n     x hi
  <chr>      <chr>      <dbl> <int> <int> <chr>
  <1> 1 Yi      C          0.1    1    3 a
  <2> 2 DJ      R          0.5    2    2 b
```


Clean names of your excel sheet

	A	B	C	D	E	F	G	H	I
1	RECORD	Age	grade --- string	race	race	bmi (kg/m^2)	weight KG		
2	931897	15 years old	10th	White	White	17.179	54.43	some random	
3	333862	17 years old	12th	White	White	20.2487	57.15	notes	
4	36253	18 years old	11th	Hispanic/Latino	Hispanic/Latino			I thought of	
5	1095530	15 years old	10th	Black or African American	Black or African American	27.9935	85.73	why	WHY?
6	1303997	14 years old	9th	All other races	Multiple - Not Hispanic/Latino	24.4922	66.68		
7	261619	17 years old	9th	All other races					
8	836640	16 years old	11th	All other races		20.5435	70.31		

```
library(readxl)
read_excel("data/messy_names.xlsx", .name_repair = janitor::make_clean_names)
```

A tibble: 20,000 x 9

	record	age	grade	string	race	race_2	bmi_kg_m_2	weight_kg	x	x_2
	<dbl>	<chr>	<chr>		<chr>	<chr>	<dbl>	<dbl>	<chr>	<chr>
1	931897	15	y...	10th	White	White	17.2	54.4	some...	<NA>
2	333862	17	y...	12th	White	White	20.2	57.2	notes	<NA>
3	36253	18		11th	White	White	NA	NA	I	<NA>

Practice

Copy and paste the code below into R to create the dataset `messy_data`

```
messy_data <- tibble(NAME = c("J N", "A C", "D E"),  
  `months follow up` = c("", 10, 11),  
  `Date of visit` = c("July 31, 2003", "Nov 12, 2005", "Aug 3, 2007"))
```

1. Clean column names with `clean_names()`.
2. Replace missing ("") data in `months_follow_up` with NA.
3. Convert `months_follow_up` to a numeric variable.
4. Convert `date_of_visit` to a date.
5. Create a column called `date_last_visit` that is the date of visit *plus* months of follow up.
6. Remove rows (cases) with missing data in `months_follow_up`.
7. Remove the spaces in `name`.

This was a *lot*, but learning R gets easier!



[Allison Horst](#)

Resources - tidyverse & data wrangling

Links

- [Learn the tidyverse](#)
- [Data wrangling cheatsheet](#)

Some of this is drawn from materials in online books/lessons:

- [R for Data Science](#) - by Garrett Grolemund & Hadley Wickham
- [Modern Dive](#) - An Introduction to Statistical and Data Sciences via R by Chester Ismay & Albert Kim
- [A gradual introduction to the tidyverse](#) - Workshop for Cascadia R 2017 by Chester Ismay and Ted Laderas
- ["Tidy Data" by Hadley Wickham](#)

Possible future workshop topics:

- reproducible reports in R (probably next) R Markdown and knitr
 - Create dynamic Word,html,pdf documents with code +output tables
- ggplot2 visualization
- advanced tidyverse: functions, purrr statistical modeling in R
-
- **Fill out feedback forms to suggest more or help us prioritize!**

Contact info:

Jessica Minnier: minnier@ohsu.edu

Meike Niederhausen: niederha@ohsu.edu

This workshop info:

- Code for these slides on github: [jminnier/berd_r_courses](https://github.com/jminnier/berd_r_courses)
- all the [R code in an R script](#)
- answers to practice problems can be found here: [html](#)